

Warp divergence

Lecture 3: control flow and synchronisation

Prof. Mike Giles

mike.giles@maths.ox.ac.uk

Oxford University Mathematical Institute

Oxford e-Research Centre

Threads are executed in warps of 32, with all threads in the warp executing the same instruction at the same time.

What happens if different threads in a warp need to do different things?

```
if (x<0.0)
    z = x-2.0;
else
    z = sqrt(x);
```

This is called *warp divergence* – CUDA will generate correct code to handle this, but to understand the performance you need to understand what CUDA does with it

Lecture 3 – p. 1

Lecture 3 – p. 2

Warp divergence

This is not a new problem.

Old CRAY vector supercomputers had a logical merge vector instruction

```
z = p ? x : y;
```

which stored the relevant element of the input vectors x, y depending on the logical vector p

```
for(i=0; i<I; i++) {
    if (p[i]) z[i] = x[i];
    else     z[i] = y[i];
}
```

Lecture 3 – p. 3

Warp divergence

Similarly, NVIDIA GPUs have *predicated* instructions which are carried out only if a logical flag is true.

```
p:  a = b + c;  // computed only if p is true
```

In the previous example, all threads compute the logical predicate and two predicated instructions

```
    p = (x<0.0);
p:  z = x-2.0;      // single instruction
!p: z = sqrt(x);
```

Lecture 3 – p. 4

Warp divergence

Note that:

- `sqrt(x)` would usually produce a NaN when `x < 0`, but it's not really executed when `x < 0` so there's no problem
- all threads execute both conditional branches, so execution cost is sum of both branches
⇒ potentially large loss of performance

Lecture 3 – p. 5

Warp divergence

Another example:

```
if (n >= 0)
    z = x[n];
else
    z = 0;
```

- `x[n]` is only read here if `n >= 0`
- don't have to worry about illegal memory accesses when `n` is negative

Lecture 3 – p. 6

Warp divergence

If the branches are big, `nvcc` compiler inserts code to check if all threads in the warp take the same branch (*warp voting*) and then branches accordingly.

```
p = ...

if (any(p)) {
p:    ...
p:    ...
}

if (any(!p)) {
!p:   ...
!p:   ...
}
```

Lecture 3 – p. 7

Warp divergence

Note:

- doesn't matter what is happening with other warps – each warp is treated separately
- if each warp only goes one way that's very efficient
- warp voting costs a few instructions, so for very simple branches the compiler just uses predication without voting

Lecture 3 – p. 8

Warp divergence

In some cases, can determine at compile time that all threads in the warp must go the same way

e.g. if `case` is a run-time argument

```
if (case==1)
    z = x*x;
else
    z = x+2.3;
```

In this case, there's no need to vote

Lecture 3 – p. 9

Warp divergence

Another example: processing a long list of elements where, depending on run-time values, a few require very expensive processing

GPU implementation:

- first process list to build two sub-lists of “simple” and “expensive” elements
- then process two sub-lists separately

Note: none of this is new – this is what we did more than 25 years ago on CRAY and Thinking Machines systems.

What's important is to understand hardware behaviour and design your algorithms / implementation accordingly

Lecture 3 – p. 11

Warp divergence

Warp divergence can lead to a big loss of parallel efficiency – one of the first things I look out for in a new application.

In worst case, effectively lose factor $32\times$ in performance if one thread needs expensive branch, while rest do nothing

Typical example: PDE application with boundary conditions

- if boundary conditions are cheap, loop over all nodes and branch as needed for boundary conditions
- if boundary conditions are expensive, use two kernels: first for interior points, second for boundary points

Lecture 3 – p. 10

Synchronisation

Already introduced `__syncthreads()`; which forms a barrier – all threads wait until every one has reached this point.

When writing conditional code, must be careful to make sure that all threads do reach the `__syncthreads()`;

Otherwise, can end up in *deadlock*

Lecture 3 – p. 12

Typical application

```
// load in data to shared memory
...
...
...

// synchronisation to ensure this has finished

__syncthreads();

// now do computation using shared data
...
...
...
```

Lecture 3 – p. 13

Warp voting

There are similar *warp voting* instructions which operate at the level of a warp:

- `int __all(predicate)`
returns non-zero (true) if all predicates in warp are true
- `int __any(predicate)`
returns non-zero (true) if any predicate is true
- `unsigned int __ballot(predicate)`
sets n^{th} bit based on n^{th} predicate

Again, I've never used these

Synchronisation

There are other synchronisation instructions which are similar but have extra capabilities:

- `int __syncthreads_count(predicate)`
counts how many predicates are true
- `int __syncthreads_and(predicate)`
returns non-zero (true) if all predicates are true
- `int __syncthreads_or(predicate)`
returns non-zero (true) if any predicate is true

I've not used these, and don't currently see a need for them

Lecture 3 – p. 14

Atomic operations

Occasionally, an application needs threads to update a counter in shared memory.

```
__shared__ int count;

...

if ( ... ) count++;
```

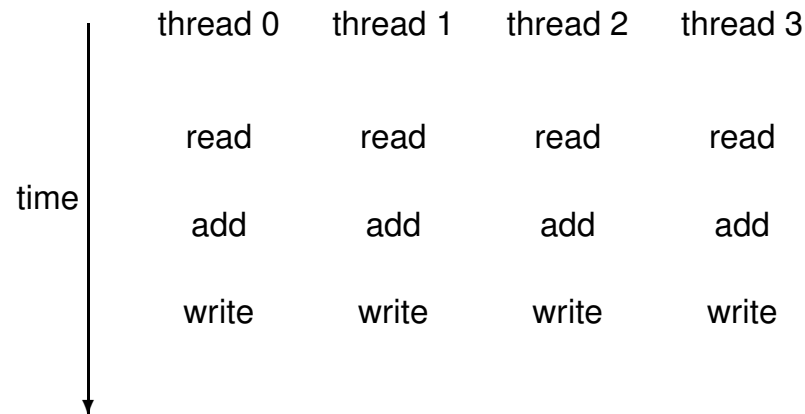
In this case, there is a problem if two (or more) threads try to do it at the same time

Lecture 3 – p. 15

Lecture 3 – p. 16

Atomic operations

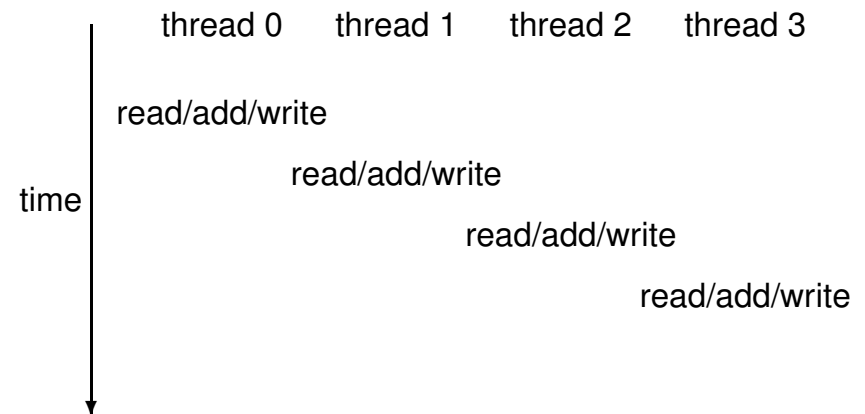
Using standard instructions, multiple threads in the same warp will only update it once.



Lecture 3 – p. 17

Atomic operations

With atomic instructions, the read/add/write becomes a single operation, and they happen one after the other



Lecture 3 – p. 18

Atomic operations

Several different atomic operations are supported, almost all only for integers:

- addition (integers, 32-bit floats – also 64-bit in Pascal)
- minimum / maximum
- increment / decrement
- exchange / compare-and-swap

These are

- not very fast for data in Kepler shared memory, better in Maxwell and Pascal
- only slightly slower for data in device global memory (operations performed in L2 cache)

Lecture 3 – p. 19

Atomic operations

Compare-and-swap:

```
int atomicCAS(int* address, int compare, int val);
```

- if `compare` equals `old` value stored at `address` then `val` is stored instead
- in either case, routine returns the value of `old`
- seems a bizarre routine at first sight, but can be very useful for atomic locks
- also can be used to implement 64-bit floating point atomic addition (now available in hardware in Pascal)

Lecture 3 – p. 20

Global atomic lock

```
// global variable: 0 unlocked, 1 locked
__device__ int lock=0;

__global__ void kernel(...) {
    ...

    if (threadIdx.x==0) {
        // set lock
        do {} while(atomicCAS(&lock,0,1));

        ...

        // free lock
        lock = 0;
    }
}
```

Lecture 3 – p. 21

Global atomic lock

Problem: when a thread writes data to device memory the order of completion is not guaranteed, so global writes may not have completed by the time the lock is unlocked

```
__global__ void kernel(...) {
    ...

    if (threadIdx.x==0) {
        do {} while(atomicCAS(&lock,0,1));
        ...
        __threadfence(); // wait for writes to finish

        // free lock
        lock = 0;
    }
}
```

Lecture 3 – p. 22

__threadfence

● `__threadfence_block();`

wait until all global and shared memory writes are visible to

- all threads in block

● `__threadfence();`

wait until all global and shared memory writes are visible to

- all threads in block
- all threads, for global data

Lecture 3 – p. 23

Atomic addition for double

```
// atomic addition from Jon Cohen at NVIDIA

static double atomicAdd(double *addr, double val)
{
    double old=*addr, assumed;

    do {
        assumed = old;
        old = __longlong_as_double(
            atomicCAS((unsigned long long int*)addr,
                __double_as_longlong(assumed),
                __double_as_longlong(val+assumed) ) );
    } while( assumed!=old );

    return old;
}
```

Lecture 3 – p. 24

Summary

- lots of esoteric capabilities – don't worry about most of them
- essential to understand warp divergence – can have a very big impact on performance
- `__syncthreads()` is vital – will see another use of it in next lecture
- the rest can be ignored until you have a critical need – then read the documentation carefully and look for examples in the SDK

Lecture 3 – p. 25

2D Laplace solver

Jacobi iteration to solve discrete Laplace equation on a uniform grid:

```
for (int j=0; j<J; j++) {
    for (int i=0; i<I; i++) {

        id = i + j*I;    // 1D memory location

        if (i==0 || i==I-1 || j==0 || j==J-1)
            u2[id] = u1[id];
        else
            u2[id] = 0.25*( u1[id-1] + u1[id+1]
                           + u1[id-I] + u1[id+I] );

    }
}
```

Lecture 3 – p. 27

Key reading

CUDA Programming Guide, version 8.0:

- Section 5.4.2: control flow and predicates
- Section 5.4.3: synchronization
- Appendix B.5: `__threadfence()` and variants
- Appendix B.6: `__syncthreads()` and variants
- Appendix B.12: atomic functions
- Appendix B.13: warp voting

Lecture 3 – p. 26


2D Laplace solver

How do we tackle this with CUDA?

- each thread responsible for one grid point
- each block of threads responsible for a block of the grid
- conceptually very similar to data partitioning in MPI distributed-memory implementations, but much simpler
- (also similar to blocking techniques to squeeze the best cache performance out of CPUs)
- great example of usefulness of 2D blocks and 2D “grid”s


Lecture 3 – p. 28

2D Laplace solver




Lecture 3 – p. 29

2D Laplace solver



Lecture 3 – p. 30

2D Laplace solver



Each block of threads processes one of these grid blocks, reading in old values and computing new values

Lecture 3 – p. 31

2D Laplace solver

```
__global__ void lap(int I, int J,
    const float* __restrict__ u1,
    float* __restrict__ u2) {

    int i = threadIdx.x + blockIdx.x*blockDim.x;
    int j = threadIdx.y + blockIdx.y*blockDim.y;
    int id = i + j*I;

    if (i==0 || i==I-1 || j==0 || j==J-1) {
        u2[id] = u1[id];    // Dirichlet b.c.'s
    }
    else {
        u2[id] = 0.25 * ( u1[id-1] + u1[id+1]
                        + u1[id-I] + u1[id+I] );
    }
}
```

Lecture 3 – p. 32

2D Laplace solver


Assumptions:

- I is a multiple of `blockDim.x`
- J is a multiple of `blockDim.y`
- hence grid breaks up perfectly into blocks

Can remove these assumptions by testing whether i, j are within grid

Lecture 3 – p. 33

2D Laplace solver



Lecture 3 – p. 34

2D Laplace solver

```
__global__ void lap(int I, int J,
                  const float* __restrict__ u1,
                  float* __restrict__ u2) {


    int i = threadIdx.x + blockIdx.x*blockDim.x;
    int j = threadIdx.y + blockIdx.y*blockDim.y;
    int id = i + j*I;

    if (i==0 || i==I-1 || j==0 || j==J-1) {
        u2[id] = u1[id]; // Dirichlet b.c.'s
    }
    else if (i<I && j<J) {
        u2[id] = 0.25f * ( u1[id-1] + u1[id+1]
                        + u1[id-I] + u1[id+I] );
    }
}
```

Lecture 3 – p. 35

2D Laplace solver

How does cache function in this application?



- if block size is a multiple of 32 in x -direction, then interior corresponds to set of complete cache lines
- “halo” points above and below are full cache lines too
- “halo” points on side are the problem – each one requires the loading of an entire cache line
- optimal block shape has aspect ratio of roughly 32:1 (or 8:1 if cache line is 32 bytes)

Lecture 3 – p. 36

3D Laplace solver

- practical 3
- each thread does an entire line in z -direction
- x, y dimensions cut up into blocks in the same way as 2D application
- `laplace3d.cu` and `laplace3d_kernel.cu` follow same approach described above
- this used to give the fastest implementation, but a new version uses 3D thread blocks, with each thread responsible for just 1 grid point
- the new version has lots more integer operations, but is still faster (due to many more active threads?)